

CSP Ex06

To run the code: install pytorch and pytorch-lightning then run the jupyter notebook. NOTE: If no gpu available set gpu=0.

Task1: See folder “data/”. It is pretty basic and works pretty well. I used L=32 and generated for each 40 temperature 4000 samples and then split in train & validation in 50-50.

Task2: I used pytorch-lightning because it is such a clean template of pytorch and very practical. Highly recommended if not known to any DL libraries. Pytorch is also the language I learned last semester at the lecture “Deep Learning”, hence I think I went a little bit overkill in this exercise xD.

Task3: My NN is a basic CNN of the form:

```
Classifier(  
    (model): Sequential(  
      (0): Conv2d(1, 6, kernel_size=(5, 5), stride=(1, 1))  
      (1): ReLU()  
      (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceiling_mode=False)  
      (3): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))  
      (4): ReLU()  
      (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceiling_mode=False)  
      (6): Flatten(start_dim=1, end_dim=-1)  
      (7): Linear(in_features=400, out_features=120, bias=True)  
      (8): ReLU()  
      (9): Linear(in_features=120, out_features=84, bias=True)  
      (10): ReLU()  
      (11): Linear(in_features=84, out_features=1, bias=True)  
      (12): Sigmoid()  
    )  
    (loss_fct): BCELoss()  
)
```

I used only one single output with a Sigmoid activation function. This is much cleaner and with this you can use the binary cross entropy loss function.

class0: 0 if $T < T_c$
class1: 1 if $T > T_c$

Task4: My binary classifier trained for like 50 epochs with Adam(lr=1e-4) and at 39 epochs it got its best result of BACC = 91.6%. I didn't neglect lattices close to the critical temperature to challenge my network. Also there are 72.5% of $T > T_c$ data points and 27.5% of $T < T_c$ due to linspace sampling in range (0.0, 8.0), where $T_c = 2.27$. Hence, the classifier is clearly more confident in classifying the samples as 0. Batch_size is set to 2048, because the bigger the better. My GPU (gtx 980) could have handled even more but I kept it at this size.

Test Error:

Accuracy: 98.4%, Avg loss: 0.038395

Balanced Accuracy: 0.9164900157868089

Classification report:

	precision	recall	f1-score	support
0.0	0.98	1.00	0.99	72323
1.0	1.00	0.83	0.91	7677
accuracy			0.98	80000
macro avg	0.99	0.92	0.95	80000
weighted avg	0.98	0.98	0.98	80000

